

# Pandas Series

```
In [1]: import pandas as pd
series = pd.Series([1,2,3,4,5])
print(series)
```

```
0    1
1    2
2    3
3    4
4    5
dtype: int64
```

## Creating a Series Using a Specified Index

```
In [2]: series = pd.Series([1,2,3,4,5], index=['a','b','c','d','c']) # note the duplicate index
print(series)
```

```
a    1
b    2
c    3
d    4
c    5
dtype: int64
```

## Accessing Elements in a Series

```
In [3]: print(series[2])          # 3
# same as
print(series.iloc[2])          # 3 - based on the position of the index
```

```
3
3
```

```
In [4]: print(series['d'])        # 4
# same as
print(series.loc['d'])          # 4 - based on the label in the index
```

```
4
4
```

```
In [5]: print(series['c'])           # more than 1 row has the index 'c'
```

```
c      3
c      5
dtype: int64
```

```
In [6]: print(series[2:])           # returns a Series
        print(series.iloc[2:])      # returns a Series
```

```
c      3
d      4
c      5
dtype: int64
c      3
d      4
c      5
dtype: int64
```

## Specifying a Datetime Range as the Index of a Series

```
In [7]: dates1 = pd.date_range('20190525', periods=12)
        print(dates1)
```

```
DatetimeIndex(['2019-05-25', '2019-05-26', '2019-05-27', '2019-05-28',
               '2019-05-29', '2019-05-30', '2019-05-31', '2019-06-01',
               '2019-06-02', '2019-06-03', '2019-06-04', '2019-06-05'],
              dtype='datetime64[ns]', freq='D')
```

```
In [8]: series = pd.Series([1,2,3,4,5,6,7,8,9,10,11,12])
        series.index = dates1
        print(series)
```

```
2019-05-25    1
2019-05-26    2
2019-05-27    3
2019-05-28    4
2019-05-29    5
2019-05-30    6
2019-05-31    7
2019-06-01    8
2019-06-02    9
2019-06-03   10
2019-06-04   11
2019-06-05   12
Freq: D, dtype: int64
```

## Date Ranges

```
In [9]: dates2 = pd.date_range('2019-05-01', periods=12, freq='M')
print(dates2)
```

```
DatetimeIndex(['2019-05-31', '2019-06-30', '2019-07-31', '2019-08-31',
               '2019-09-30', '2019-10-31', '2019-11-30', '2019-12-31',
               '2020-01-31', '2020-02-29', '2020-03-31', '2020-04-30'],
              dtype='datetime64[ns]', freq='M')
```

```
In [10]: dates2 = pd.date_range('2019-05-01', periods=12, freq='MS')
print(dates2)
```

```
DatetimeIndex(['2019-05-01', '2019-06-01', '2019-07-01', '2019-08-01',
               '2019-09-01', '2019-10-01', '2019-11-01', '2019-12-01',
               '2020-01-01', '2020-02-01', '2020-03-01', '2020-04-01'],
              dtype='datetime64[ns]', freq='MS')
```

```
In [11]: dates2 = pd.date_range('05-01-2019', periods=12, freq='MS')
print(dates2)
```

```
DatetimeIndex(['2019-05-01', '2019-06-01', '2019-07-01', '2019-08-01',
               '2019-09-01', '2019-10-01', '2019-11-01', '2019-12-01',
               '2020-01-01', '2020-02-01', '2020-03-01', '2020-04-01'],
              dtype='datetime64[ns]', freq='MS')
```

```
In [12]: dates3 = pd.date_range('2019/05/17 09:00:00', periods=8, freq='H')
print(dates3)
```

```
DatetimeIndex(['2019-05-17 09:00:00', '2019-05-17 10:00:00',
               '2019-05-17 11:00:00', '2019-05-17 12:00:00',
               '2019-05-17 13:00:00', '2019-05-17 14:00:00',
               '2019-05-17 15:00:00', '2019-05-17 16:00:00'],
              dtype='datetime64[ns]', freq='H')
```

## Pandas DataFrame

```
In [5]: import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(5,4),
                  columns=list('ABCD'))
print(df)
```

```
      A         B         C         D
0  1.176839 -0.568883  0.175213 -0.702281
1 -1.719396  0.408782  0.856516 -0.480629
2  0.274366  0.868352 -0.404123 -1.682070
3 -1.965048  0.321428 -1.259396 -1.029488
4  0.802132  1.194193 -0.502426  0.393425
```

```
In [4]: np.random.randn?
```

## Specifying the Index in a DataFrame

```
In [14]: df = pd.read_csv('data.csv') # load dataframe from CSV file
days = pd.date_range('20190525', periods=10)
df.index = days
print(df)
```

	A	B	C	D
2019-05-25	0.187497	1.122150	-0.988277	-1.985934
2019-05-26	0.360803	-0.562243	-0.340693	-0.986988
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-28	-0.279572	-0.702492	0.252265	0.958977
2019-05-29	0.537438	-1.737568	0.714727	-0.939288
2019-05-30	0.070011	-0.516443	-1.655689	0.246721
2019-05-31	0.001268	0.951517	2.107360	-0.108726
2019-06-01	-0.185258	0.856520	-0.686285	1.104195
2019-06-02	0.387023	1.706336	-2.452653	0.260466
2019-06-03	-1.054974	0.556775	-0.945219	-0.030295

```
In [15]: print(df.index)
```

```
DatetimeIndex(['2019-05-25', '2019-05-26', '2019-05-27', '2019-05-28',
               '2019-05-29', '2019-05-30', '2019-05-31', '2019-06-01',
               '2019-06-02', '2019-06-03'],
              dtype='datetime64[ns]', freq='D')
```

```
In [16]: print(df.values)
```

```
[[ 1.874970e-01  1.122150e+00 -9.882770e-01 -1.985934e+00]
 [ 3.608030e-01 -5.622430e-01 -3.406930e-01 -9.869880e-01]
 [-4.062700e-02  6.733300e-02 -4.529780e-01  6.862230e-01]
 [-2.795720e-01 -7.024920e-01  2.522650e-01  9.589770e-01]
 [ 5.374380e-01 -1.737568e+00  7.147270e-01 -9.392880e-01]
 [ 7.001100e-02 -5.164430e-01 -1.655689e+00  2.467210e-01]
 [ 1.268000e-03  9.515170e-01  2.107360e+00 -1.087260e-01]
 [-1.852580e-01  8.565200e-01 -6.862850e-01  1.104195e+00]
 [ 3.870230e-01  1.706336e+00 -2.452653e+00  2.604660e-01]
 [-1.054974e+00  5.567750e-01 -9.452190e-01 -3.029500e-02]]
```

```
In [17]: print(df.describe())
```

	A	B	C	D
count	10.000000	10.000000	10.000000	10.000000
mean	-0.001639	0.174188	-0.444744	-0.079465
std	0.451656	1.049677	1.267397	0.971164
min	-1.054974	-1.737568	-2.452653	-1.985934
25%	-0.149100	-0.550793	-0.977513	-0.731647
50%	0.035640	0.312054	-0.569632	0.108213
75%	0.317477	0.927768	0.104026	0.579784
max	0.537438	1.706336	2.107360	1.104195

```
In [18]: print(df.mean(0))    # 0 means compute the mean for each columns
```

```
A    -0.001639
B     0.174188
C    -0.444744
D    -0.079465
dtype: float64
```

```
In [19]: print(df.mean(1))    # 1 means compute the mean for each row
```

```
2019-05-25    -0.416141
2019-05-26    -0.382280
2019-05-27     0.064988
2019-05-28     0.057294
2019-05-29    -0.356173
2019-05-30    -0.463850
2019-05-31     0.737855
2019-06-01     0.272293
2019-06-02    -0.024707
2019-06-03    -0.368428
Freq: D, dtype: float64
```

## Extracting from DataFrames

### Selecting the First and Last Five Rows

```
In [20]: print(df.head())
```

	A	B	C	D
2019-05-25	0.187497	1.122150	-0.988277	-1.985934
2019-05-26	0.360803	-0.562243	-0.340693	-0.986988
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-28	-0.279572	-0.702492	0.252265	0.958977
2019-05-29	0.537438	-1.737568	0.714727	-0.939288

```
In [21]: print(df.head(8))      # prints out the first 8 rows
```

	A	B	C	D
2019-05-25	0.187497	1.122150	-0.988277	-1.985934
2019-05-26	0.360803	-0.562243	-0.340693	-0.986988
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-28	-0.279572	-0.702492	0.252265	0.958977
2019-05-29	0.537438	-1.737568	0.714727	-0.939288
2019-05-30	0.070011	-0.516443	-1.655689	0.246721
2019-05-31	0.001268	0.951517	2.107360	-0.108726
2019-06-01	-0.185258	0.856520	-0.686285	1.104195

```
In [22]: print(df.tail())
```

	A	B	C	D
2019-05-30	0.070011	-0.516443	-1.655689	0.246721
2019-05-31	0.001268	0.951517	2.107360	-0.108726
2019-06-01	-0.185258	0.856520	-0.686285	1.104195
2019-06-02	0.387023	1.706336	-2.452653	0.260466
2019-06-03	-1.054974	0.556775	-0.945219	-0.030295

```
In [23]: print(df.tail(8))      # prints out the last 8 rows
```

	A	B	C	D
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-28	-0.279572	-0.702492	0.252265	0.958977
2019-05-29	0.537438	-1.737568	0.714727	-0.939288
2019-05-30	0.070011	-0.516443	-1.655689	0.246721
2019-05-31	0.001268	0.951517	2.107360	-0.108726
2019-06-01	-0.185258	0.856520	-0.686285	1.104195
2019-06-02	0.387023	1.706336	-2.452653	0.260466
2019-06-03	-1.054974	0.556775	-0.945219	-0.030295

## Selecting a Specific Column in a DataFrame

```
In [24]: print(df['A'])
# same as
print(df.A)
```

```
2019-05-25    0.187497
2019-05-26    0.360803
2019-05-27   -0.040627
2019-05-28   -0.279572
2019-05-29    0.537438
2019-05-30    0.070011
2019-05-31    0.001268
2019-06-01   -0.185258
2019-06-02    0.387023
2019-06-03   -1.054974
Freq: D, Name: A, dtype: float64
2019-05-25    0.187497
2019-05-26    0.360803
2019-05-27   -0.040627
2019-05-28   -0.279572
2019-05-29    0.537438
2019-05-30    0.070011
2019-05-31    0.001268
2019-06-01   -0.185258
2019-06-02    0.387023
2019-06-03   -1.054974
Freq: D, Name: A, dtype: float64
```

```
In [25]: print(df[['A', 'B']])
```

```
           A         B
2019-05-25  0.187497  1.122150
2019-05-26  0.360803 -0.562243
2019-05-27 -0.040627  0.067333
2019-05-28 -0.279572 -0.702492
2019-05-29  0.537438 -1.737568
2019-05-30  0.070011 -0.516443
2019-05-31  0.001268  0.951517
2019-06-01 -0.185258  0.856520
2019-06-02  0.387023  1.706336
2019-06-03 -1.054974  0.556775
```

## Slicing Based on Row Number

```
In [26]: print(df[2:4])
```

```
           A         B         C         D
2019-05-27 -0.040627  0.067333 -0.452978  0.686223
2019-05-28 -0.279572 -0.702492  0.252265  0.958977
```

```
In [27]: print(df.iloc[2:4])      # 2 rows
```

	A	B	C	D
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-28	-0.279572	-0.702492	0.252265	0.958977

```
In [28]: print(df.iloc[2:5])      # 3 rows
```

	A	B	C	D
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-28	-0.279572	-0.702492	0.252265	0.958977
2019-05-29	0.537438	-1.737568	0.714727	-0.939288

```
In [29]: print(df.iloc[[2,4]])      # 2 rows
```

	A	B	C	D
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-29	0.537438	-1.737568	0.714727	-0.939288

```
In [30]: # print(df[[2,4]])      # error; need to use the iloc indexer
print(df.iloc[2])      # prints out row number 2
```

```
A    -0.040627
B     0.067333
C    -0.452978
D     0.686223
Name: 2019-05-27 00:00:00, dtype: float64
```

## Slicing Based on Row and Column Numbers

```
In [31]: print(df.iloc[2:4, 1:4])      # 2 rows, 3 columns
```

	B	C	D
2019-05-27	0.067333	-0.452978	0.686223
2019-05-28	-0.702492	0.252265	0.958977

```
In [32]: print(df.iloc[[2,4], [1,3]])      # 2 rows, 2 columns
```

	B	D
2019-05-27	0.067333	0.686223
2019-05-29	-1.737568	-0.939288

## Slicing Based on Labels



```
In [33]: print(df['20190601':'20190603'])
```

	A	B	C	D
2019-06-01	-0.185258	0.856520	-0.686285	1.104195
2019-06-02	0.387023	1.706336	-2.452653	0.260466
2019-06-03	-1.054974	0.556775	-0.945219	-0.030295

```
In [34]: print(df.loc['20190601':'20190603'])
```

	A	B	C	D
2019-06-01	-0.185258	0.856520	-0.686285	1.104195
2019-06-02	0.387023	1.706336	-2.452653	0.260466
2019-06-03	-1.054974	0.556775	-0.945219	-0.030295

```
In [35]: print(df.loc['20190601':'20190603', 'A':'C'])
```

	A	B	C
2019-06-01	-0.185258	0.856520	-0.686285
2019-06-02	0.387023	1.706336	-2.452653
2019-06-03	-1.054974	0.556775	-0.945219

```
In [36]: print(df.loc['20190601':'20190603', ['A', 'C']])
```

	A	C
2019-06-01	-0.185258	-0.686285
2019-06-02	0.387023	-2.452653
2019-06-03	-1.054974	-0.945219

```
In [37]: print(df.loc['20190601'])
```

```
A    -0.185258
B     0.856520
C    -0.686285
D     1.104195
Name: 2019-06-01 00:00:00, dtype: float64
```

```
In [38]: # print(df.loc[['20190601', '20190603']]) # KeyError
```

```
In [12]: from datetime import datetime
date1 = datetime(2019, 6, 1, 0, 0, 0)
date2 = datetime(2019, 6, 3, 0, 0, 0)
print(df.loc[[date1,date2]])
```

```
-----
--
KeyError                                Traceback (most recent call last)
<ipython-input-12-b7636d7fc292> in <module>
      2 date1 = datetime(2019, 6, 1, 0, 0, 0)
      3 date2 = datetime(2019, 6, 3, 0, 0, 0)
----> 4 print(df.loc[[date1,date2]])

~/anaconda3/lib/python3.7/site-packages/pandas/core/indexing.py in __getitem__(self, key)
    1498
    1499         maybe_callable = com.apply_if_callable(key, self.obj)
-> 1500         return self._getitem_axis(maybe_callable, axis=axis)
    1501
    1502     def _is_scalar_access(self, key):

~/anaconda3/lib/python3.7/site-packages/pandas/core/indexing.py in _getitem_axis(self, key, axis)
    1900         raise ValueError('Cannot index with multidimensional key')
    1901
-> 1902         return self._getitem_iterable(key, axis=axis)
    1903
    1904         # nested tuple slicing

~/anaconda3/lib/python3.7/site-packages/pandas/core/indexing.py in _getitem_iterable(self, key, axis)
    1203         # A collection of keys
    1204         keyarr, indexer = self._get_listlike_indexer(key, axis,
s,
-> 1205                                                     raise_missing=
ssing=False)
    1206         return self.obj._reindex_with_indexers({axis: [keyarr,
r, indexer]},
    1207                                                     copy=True, all
ow_dups=True)

~/anaconda3/lib/python3.7/site-packages/pandas/core/indexing.py in _get_listlike_indexer(self, key, axis, raise_missing)
    1159         self._validate_read_indexer(keyarr, indexer,
    1160                                     o._get_axis_number(axis),
-> 1161                                     raise_missing=raise_missing)
    1162         return keyarr, indexer
    1163

~/anaconda3/lib/python3.7/site-packages/pandas/core/indexing.py in _validate_read_indexer(self, key, indexer, axis, raise_missing)
    1244         raise KeyError(
    1245             u"None of [{key}] are in the [{axis}]"
.
format
(
```

```
-> 1246                                     key=key, axis=self.obj._get_axis_name(axi
s)))
    1247
    1248                                     # We (temporarily) allow for some missing keys with .
loc, except in

KeyError: "None of [DatetimeIndex(['2019-06-01', '2019-06-03'], dtype='da
atetime64[ns]', freq=None)] are in the [index]"
```

```
In [40]: print(df.loc[date1, ['A', 'C']])
```

```
A    -0.185258
C    -0.686285
Name: 2019-06-01 00:00:00, dtype: float64
```

## Selecting a Single Cell in a DataFrame

```
In [14]: from datetime import datetime
d = datetime(2019, 6, 3, 0, 0, 0)
print(df.at[d, 'B'])
```

```
-----
--
ValueError                                Traceback (most recent call last)
<ipython-input-14-2df6b65fb828> in <module>
      1 from datetime import datetime
      2 d = datetime(2019, 6, 3, 0, 0, 0)
----> 3 print(df.at[d, 'B'])

~/anaconda3/lib/python3.7/site-packages/pandas/core/indexing.py in __getitem__(self, key)
    2267         raise ValueError('Invalid call for scalar access (getting)!')
    2268
-> 2269         key = self._convert_key(key)
    2270         return self.obj._get_value(*key, takeable=self._takeable)
    2271

~/anaconda3/lib/python3.7/site-packages/pandas/core/indexing.py in _convert_key(self, key, is_setter)
    2349         if ax.is_integer():
    2350             if not is_integer(i):
-> 2351                 raise ValueError("At based indexing on an integer index "
    2352                                 "can only have integer indexers")
    2353         else:
```

**ValueError:** At based indexing on an integer index can only have integer indexers

## Selecting Based on Cell Value

```
In [42]: print(df[(df.A > 0) & (df.B>0)])
```

	A	B	C	D
2019-05-25	0.187497	1.122150	-0.988277	-1.985934
2019-05-31	0.001268	0.951517	2.107360	-0.108726
2019-06-02	0.387023	1.706336	-2.452653	0.260466

## Transforming DataFrames

```
In [43]: print(df.transpose())
```

	2019-05-25	2019-05-26	2019-05-27	2019-05-28	2019-05-29	2019-05-30
\						
A	0.187497	0.360803	-0.040627	-0.279572	0.537438	0.070011
B	1.122150	-0.562243	0.067333	-0.702492	-1.737568	-0.516443
C	-0.988277	-0.340693	-0.452978	0.252265	0.714727	-1.655689
D	-1.985934	-0.986988	0.686223	0.958977	-0.939288	0.246721

  

	2019-05-31	2019-06-01	2019-06-02	2019-06-03
A	0.001268	-0.185258	0.387023	-1.054974
B	0.951517	0.856520	1.706336	0.556775
C	2.107360	-0.686285	-2.452653	-0.945219
D	-0.108726	1.104195	0.260466	-0.030295

```
In [44]: print(df.T)
```

	2019-05-25	2019-05-26	2019-05-27	2019-05-28	2019-05-29	2019-05-30
\						
A	0.187497	0.360803	-0.040627	-0.279572	0.537438	0.070011
B	1.122150	-0.562243	0.067333	-0.702492	-1.737568	-0.516443
C	-0.988277	-0.340693	-0.452978	0.252265	0.714727	-1.655689
D	-1.985934	-0.986988	0.686223	0.958977	-0.939288	0.246721

  

	2019-05-31	2019-06-01	2019-06-02	2019-06-03
A	0.001268	-0.185258	0.387023	-1.054974
B	0.951517	0.856520	1.706336	0.556775
C	2.107360	-0.686285	-2.452653	-0.945219
D	-0.108726	1.104195	0.260466	-0.030295

```
In [45]: def checkSeriesOrDataframe(var):
    if isinstance(var, pd.DataFrame):
        return 'Dataframe'
    if isinstance(var, pd.Series):
        return 'Series'
```

## Sorting Data in a DataFrame

### Sorting by Index

```
In [46]: print(df.sort_index(axis=0, ascending=False)) # axis = 0 means sort by
# index
```

	A	B	C	D
2019-06-03	-1.054974	0.556775	-0.945219	-0.030295
2019-06-02	0.387023	1.706336	-2.452653	0.260466
2019-06-01	-0.185258	0.856520	-0.686285	1.104195
2019-05-31	0.001268	0.951517	2.107360	-0.108726
2019-05-30	0.070011	-0.516443	-1.655689	0.246721
2019-05-29	0.537438	-1.737568	0.714727	-0.939288
2019-05-28	-0.279572	-0.702492	0.252265	0.958977
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-26	0.360803	-0.562243	-0.340693	-0.986988
2019-05-25	0.187497	1.122150	-0.988277	-1.985934

```
In [47]: print(df.sort_index(axis=1, ascending=False)) # axis = 1 means sort by
# column
```

	D	C	B	A
2019-05-25	-1.985934	-0.988277	1.122150	0.187497
2019-05-26	-0.986988	-0.340693	-0.562243	0.360803
2019-05-27	0.686223	-0.452978	0.067333	-0.040627
2019-05-28	0.958977	0.252265	-0.702492	-0.279572
2019-05-29	-0.939288	0.714727	-1.737568	0.537438
2019-05-30	0.246721	-1.655689	-0.516443	0.070011
2019-05-31	-0.108726	2.107360	0.951517	0.001268
2019-06-01	1.104195	-0.686285	0.856520	-0.185258
2019-06-02	0.260466	-2.452653	1.706336	0.387023
2019-06-03	-0.030295	-0.945219	0.556775	-1.054974

## Sorting by Value

```
In [48]: print(df.sort_values('A', axis=0))
```

	A	B	C	D
2019-06-03	-1.054974	0.556775	-0.945219	-0.030295
2019-05-28	-0.279572	-0.702492	0.252265	0.958977
2019-06-01	-0.185258	0.856520	-0.686285	1.104195
2019-05-27	-0.040627	0.067333	-0.452978	0.686223
2019-05-31	0.001268	0.951517	2.107360	-0.108726
2019-05-30	0.070011	-0.516443	-1.655689	0.246721
2019-05-25	0.187497	1.122150	-0.988277	-1.985934
2019-05-26	0.360803	-0.562243	-0.340693	-0.986988
2019-06-02	0.387023	1.706336	-2.452653	0.260466
2019-05-29	0.537438	-1.737568	0.714727	-0.939288

```
In [49]: print(df.sort_values('20190601', axis=1))
```

	C	A	B	D
2019-05-25	-0.988277	0.187497	1.122150	-1.985934
2019-05-26	-0.340693	0.360803	-0.562243	-0.986988
2019-05-27	-0.452978	-0.040627	0.067333	0.686223
2019-05-28	0.252265	-0.279572	-0.702492	0.958977
2019-05-29	0.714727	0.537438	-1.737568	-0.939288
2019-05-30	-1.655689	0.070011	-0.516443	0.246721
2019-05-31	2.107360	0.001268	0.951517	-0.108726
2019-06-01	-0.686285	-0.185258	0.856520	1.104195
2019-06-02	-2.452653	0.387023	1.706336	0.260466
2019-06-03	-0.945219	-1.054974	0.556775	-0.030295

## Applying Functions to a DataFrame

```
In [50]: import math
sq_root = lambda x: math.sqrt(x) if x > 0 else x
sq       = lambda x: x**2
```

```
In [51]: print(df.B.apply(sq_root))
```

2019-05-25	1.059316
2019-05-26	-0.562243
2019-05-27	0.259486
2019-05-28	-0.702492
2019-05-29	-1.737568
2019-05-30	-0.516443
2019-05-31	0.975457
2019-06-01	0.925484
2019-06-02	1.306268
2019-06-03	0.746174

Freq: D, Name: B, dtype: float64

```
In [52]: print(df.B.apply(sq))
```

2019-05-25	1.259221
2019-05-26	0.316117
2019-05-27	0.004534
2019-05-28	0.493495
2019-05-29	3.019143
2019-05-30	0.266713
2019-05-31	0.905385
2019-06-01	0.733627
2019-06-02	2.911583
2019-06-03	0.309998

Freq: D, Name: B, dtype: float64

```
In [53]: # df.apply(sq_root)      # ValueError
```

```
In [54]: df.apply(sq)
```

```
Out[54]:
```

	A	B	C	D
2019-05-25	0.035155	1.259221	0.976691	3.943934
2019-05-26	0.130179	0.316117	0.116072	0.974145
2019-05-27	0.001651	0.004534	0.205189	0.470902
2019-05-28	0.078161	0.493495	0.063638	0.919637
2019-05-29	0.288840	3.019143	0.510835	0.882262
2019-05-30	0.004902	0.266713	2.741306	0.060871
2019-05-31	0.000002	0.905385	4.440966	0.011821
2019-06-01	0.034321	0.733627	0.470987	1.219247
2019-06-02	0.149787	2.911583	6.015507	0.067843
2019-06-03	1.112970	0.309998	0.893439	0.000918

```
In [55]: for column in df:
          df[column] = df[column].apply(sq_root)
          print(df)
```

	A	B	C	D
2019-05-25	0.433009	1.059316	-0.988277	-1.985934
2019-05-26	0.600669	-0.562243	-0.340693	-0.986988
2019-05-27	-0.040627	0.259486	-0.452978	0.828386
2019-05-28	-0.279572	-0.702492	0.502260	0.979274
2019-05-29	0.733102	-1.737568	0.845415	-0.939288
2019-05-30	0.264596	-0.516443	-1.655689	0.496710
2019-05-31	0.035609	0.975457	1.451675	-0.108726
2019-06-01	-0.185258	0.925484	-0.686285	1.050807
2019-06-02	0.622112	1.306268	-2.452653	0.510359
2019-06-03	-1.054974	0.746174	-0.945219	-0.030295

```
In [56]: print(df.apply(np.sum, axis=0))
```

```
A    1.128665
B    1.753438
C   -4.722444
D   -0.185696
dtype: float64
```



```
In [57]: print(df.apply(np.sum, axis=1))
```

```
2019-05-25    -1.481886
2019-05-26    -1.289255
2019-05-27     0.594267
2019-05-28     0.499470
2019-05-29    -1.098339
2019-05-30    -1.410826
2019-05-31     2.354015
2019-06-01     1.104747
2019-06-02    -0.013915
2019-06-03    -1.284314
Freq: D, dtype: float64
```

## Adding and Removing Rows and Columns in a DataFrame

```
In [58]: import pandas as pd
```

```
data = {'name': ['Janet', 'Nad', 'Timothy', 'June', 'Amy'],
        'year': [2012, 2012, 2013, 2014, 2014],
        'reports': [6, 13, 14, 1, 7]}

df = pd.DataFrame(data, index =
    ['Singapore', 'China', 'Japan', 'Sweden', 'Norway'])
print(df)
```

	name	year	reports
Singapore	Janet	2012	6
China	Nad	2012	13
Japan	Timothy	2013	14
Sweden	June	2014	1
Norway	Amy	2014	7

## Adding a Column

```
In [59]: import numpy as np
```

```
schools = np.array(["Cambridge", "Oxford", "Oxford", "Cambridge", "Oxford"])
df["school"] = schools
print(df)
```

	name	year	reports	school
Singapore	Janet	2012	6	Cambridge
China	Nad	2012	13	Oxford
Japan	Timothy	2013	14	Oxford
Sweden	June	2014	1	Cambridge
Norway	Amy	2014	7	Oxford

## Removing Rows

```
In [60]: print(df.drop(['China', 'Japan'])) # drop rows based on value of index
```

	name	year	reports	school
Singapore	Janet	2012	6	Cambridge
Sweden	June	2014	1	Cambridge
Norway	Amy	2014	7	Oxford

```
In [61]: print(df[df.name != 'Nad']) # drop row based on column value
```

	name	year	reports	school
Singapore	Janet	2012	6	Cambridge
Japan	Timothy	2013	14	Oxford
Sweden	June	2014	1	Cambridge
Norway	Amy	2014	7	Oxford

```
In [62]: print(df.drop(df.index[1]))
# same as df.drop['China']
```

	name	year	reports	school
Singapore	Janet	2012	6	Cambridge
Japan	Timothy	2013	14	Oxford
Sweden	June	2014	1	Cambridge
Norway	Amy	2014	7	Oxford

```
In [63]: print(df.drop(df.index[[1,2]])) # remove the second and third row
```

	name	year	reports	school
Singapore	Janet	2012	6	Cambridge
Sweden	June	2014	1	Cambridge
Norway	Amy	2014	7	Oxford

```
In [64]: print(df.drop(df.index[-2])) # remove second last row
```

	name	year	reports	school
Singapore	Janet	2012	6	Cambridge
China	Nad	2012	13	Oxford
Japan	Timothy	2013	14	Oxford
Norway	Amy	2014	7	Oxford

## Removing Columns

```
In [65]: print(df.drop('reports', axis=1))    # drop column
```

	name	year	school
Singapore	Janet	2012	Cambridge
China	Nad	2012	Oxford
Japan	Timothy	2013	Oxford
Sweden	June	2014	Cambridge
Norway	Amy	2014	Oxford

```
In [66]: print(df.drop(df.columns[1], axis=1))    # drop using columns number
```

	name	reports	school
Singapore	Janet	6	Cambridge
China	Nad	13	Oxford
Japan	Timothy	14	Oxford
Sweden	June	1	Cambridge
Norway	Amy	7	Oxford

```
In [67]: print(df.drop(df.columns[[1,3]], axis=1))    # drop multiple columns
```

	name	reports
Singapore	Janet	6
China	Nad	13
Japan	Timothy	14
Sweden	June	1
Norway	Amy	7

## Generating a Crosstab

```
In [68]: df = pd.DataFrame(
    {
        "Gender": ['Male', 'Male', 'Female', 'Female', 'Female'],
        "Team"   : [1,2,3,3,1]
    })
print(df)
```

	Gender	Team
0	Male	1
1	Male	2
2	Female	3
3	Female	3
4	Female	1

```
In [69]: print("Displaying the distribution of genders in each team")  
print(pd.crosstab(df.Gender, df.Team))
```

```
Displaying the distribution of genders in each team  
Team    1  2  3  
Gender  
Female   1  0  2  
Male     1  1  0
```

```
In [70]: print(pd.crosstab(df.Team, df.Gender))
```

```
Gender  Female  Male  
Team  
1         1     1  
2         0     1  
3         2     0
```

```
In [ ]:
```