

# Getting Datasets

## Using the Scikit-learn Dataset

```
In [1]: from sklearn import datasets  
iris = datasets.load_iris()    # raw data of type Bunch
```

```
In [2]: print(iris.DESCR)
```

```
Iris Plants Database
```

```
=====
```

```
Notes
```

```
-----
```

```
Data Set Characteristics:
```

```
:Number of Instances: 150 (50 in each of three classes)
:Number of Attributes: 4 numeric, predictive attributes and the class
:Attribute Information:
  - sepal length in cm
  - sepal width in cm
  - petal length in cm
  - petal width in cm
  - class:
    - Iris-Setosa
    - Iris-Versicolour
    - Iris-Virginica
```

```
:Summary Statistics:
```

```
=====  =====  =====  =====  =====  =====
              Min    Max    Mean    SD    Class Correlation
=====  =====  =====  =====  =====  =====
sepal length:  4.3    7.9    5.84    0.83    0.7826
sepal width:   2.0    4.4    3.05    0.43   -0.4194
petal length:  1.0    6.9    3.76    1.76    0.9490 (high!)
petal width:   0.1    2.5    1.20    0.76    0.9565 (high!)
=====  =====  =====  =====  =====  =====
```

```
:Missing Attribute Values: None
:Class Distribution: 33.3% for each of 3 classes.
:Creator: R.A. Fisher
:Donor: Michael Marshall (MARSHALL%PLU@io.arc.nasa.gov)
:Date: July, 1988
```

This is a copy of UCI ML iris datasets.

<http://archive.ics.uci.edu/ml/datasets/Iris> (<http://archive.ics.uci.edu/ml/datasets/Iris>)

The famous Iris database, first used by Sir R.A Fisher

This is perhaps the best known database to be found in the pattern recognition literature. Fisher's paper is a classic in the field and is referenced frequently to this day. (See Duda & Hart, for example.) The data set contains 3 classes of 50 instances each, where each class refers to a type of iris plant. One class is linearly separable from the other 2; the latter are NOT linearly separable from each other.

```
References
```

```
-----
```

```
- Fisher,R.A. "The use of multiple measurements in taxonomic problems"
```

Annual Eugenics, 7, Part II, 179-188 (1936); also in "Contributions to Mathematical Statistics" (John Wiley, NY, 1950).

- Duda, R.O., & Hart, P.E. (1973) Pattern Classification and Scene Analysis. (Q327.D83) John Wiley & Sons. ISBN 0-471-22361-1. See page 218.
- Dasarathy, B.V. (1980) "Nosing Around the Neighborhood: A New System Structure and Classification Rule for Recognition in Partially Exposed Environments". IEEE Transactions on Pattern Analysis and Machine Intelligence, Vol. PAMI-2, No. 1, 67-71.
- Gates, G.W. (1972) "The Reduced Nearest Neighbor Rule". IEEE Transactions on Information Theory, May 1972, 431-433.
- See also: 1988 MLC Proceedings, 54-64. Cheeseman et al's AUTOCLASS II conceptual clustering system finds 3 classes in the data.
- Many, many more ...

```
In [3]: print(iris.data)           # Features
```

```
[[5.1 3.5 1.4 0.2]
 [4.9 3.  1.4 0.2]
 [4.7 3.2 1.3 0.2]
 [4.6 3.1 1.5 0.2]
 [5.  3.6 1.4 0.2]
 [5.4 3.9 1.7 0.4]
 [4.6 3.4 1.4 0.3]
 [5.  3.4 1.5 0.2]
 [4.4 2.9 1.4 0.2]
 [4.9 3.1 1.5 0.1]
 [5.4 3.7 1.5 0.2]
 [4.8 3.4 1.6 0.2]
 [4.8 3.  1.4 0.1]
 [4.3 3.  1.1 0.1]
 [5.8 4.  1.2 0.2]
 [5.7 4.4 1.5 0.4]
 [5.4 3.9 1.3 0.4]
 [5.1 3.5 1.4 0.3]
 [5.7 3.8 1.7 0.3]
 [5.  3.  1.  0.  ]]
```

```
In [4]: print(iris.feature_names)  # Feature Names
```


```
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
```



```
In [8]: %matplotlib inline
from matplotlib import pyplot as plt
from sklearn.datasets.samples_generator import make_regression

X, y = make_regression(n_samples=100, n_features=1, noise=5.4)
plt.scatter(X,y)
```

Out[8]: <matplotlib.collections.PathCollection at 0x1a233d4dd8>



## Clustered Dataset


```
In [9]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
from sklearn.datasets import make_blobs

X, y = make_blobs(500, centers=3) # Generate isotropic Gaussian
                                   # blobs for clustering

rgb = np.array(['r', 'g', 'b'])

# plot the blobs using a scatter plot and use color coding
plt.scatter(X[:, 0], X[:, 1], color=rgb[y])
```

Out[9]: <matplotlib.collections.PathCollection at 0x1a234e9f98>




## Clustered Dataset Distributed in Circular Fashion

```
In [10]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
from sklearn.datasets import make_circles

X, y = make_circles(n_samples=100, noise=0.09)

rgb = np.array(['r', 'g', 'b'])
plt.scatter(X[:, 0], X[:, 1], color=rgb[y])
```

Out[10]: <matplotlib.collections.PathCollection at 0x1a235b25c0>



## Getting Started with Scikit-learn

```
In [11]: %matplotlib inline
import matplotlib.pyplot as plt


# represents the heights of a group of people in metres
heights = [[1.6], [1.65], [1.7], [1.73], [1.8]]

# represents the weights of a group of people in kgs
weights = [[60], [65], [72.3], [75], [80]]

plt.title('Weights plotted against heights')
plt.xlabel('Heights in metres')
plt.ylabel('Weights in kilograms')

plt.plot(heights, weights, 'k.')

# axis range for x and y
plt.axis([1.5, 1.85, 50, 90])
plt.grid(True)
```



## Using the LinearRegression Class for Fitting the Model

```
In [12]: from sklearn.linear_model import LinearRegression

# Create and fit the model
model = LinearRegression()
model.fit(X=heights, y=weights)
```

```
Out[12]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

## Making Predictions



```
In [13]: # make prediction
weight = model.predict([[1.75]])[0][0]
print(round(weight,2))           # 76.04
```

76.04

## Plotting the Linear Regression Line

```
In [14]: import matplotlib.pyplot as plt

heights = [[1.6], [1.65], [1.7], [1.73], [1.8]]
weights = [[60], [65], [72.3], [75], [80]]


plt.title('Weights plotted against heights')
plt.xlabel('Heights in metres')
plt.ylabel('Weights in kilograms')

plt.plot(heights, weights, 'k.')

plt.axis([1.5, 1.85, 50, 90])
plt.grid(True)

# plot the regression line
plt.plot(heights, model.predict(heights), color='r')
```

Out[14]: [<matplotlib.lines.Line2D at 0x1a23b640b8>]



## Getting the Gradient and Intercept of the Linear Regression Line


```
In [15]: plt.title('Weights plotted against heights')
plt.xlabel('Heights in metres')
plt.ylabel('Weights in kilograms')

plt.plot(heights, weights, 'k.')

plt.axis([0, 1.85, -200, 200])
plt.grid(True)

# plot the regression line
extreme_heights = [[0], [1.8]]
plt.plot(extreme_heights, model.predict(extreme_heights), color='b')
```

Out[15]: [ <matplotlib.lines.Line2D at 0x1a23c43f98>]



```
In [16]: round(model.predict([[0]])[0][0],2)    # -104.75
```

Out[16]: -104.75

```
In [17]: print(round(model.intercept_[0],2))    # -104.75
```

-104.75

```
In [18]: print(round(model.coef_[0][0],2))     # 103.31
```

103.31

## Examining the Performance of the Model by Calculating the Residual Sum of Squares

```
In [19]: import numpy as np

print('Residual sum of squares: %.2f' %
      np.sum((weights - model.predict(heights)) ** 2))
```

Residual sum of squares: 5.34

The RSS should be as small as possible, with 0 indicating that the regression line fits the points exactly (rarely achievable in the real world).

## Evaluating the Model Using a Test Dataset

```
In [20]: # test data
heights_test = [[1.58], [1.62], [1.69], [1.76], [1.82]]
weights_test = [[58], [63], [72], [73], [85]]
```

```
In [21]: # Total Sum of Squares (TSS)
weights_test_mean = np.mean(np.ravel(weights_test))
TSS = np.sum((np.ravel(weights_test) -
              weights_test_mean) ** 2)
print("TSS: %.2f" % TSS)

# Residual Sum of Squares (RSS)
RSS = np.sum((np.ravel(weights_test) -
              np.ravel(model.predict(heights_test)))
              ** 2)
print("RSS: %.2f" % RSS)

# R_squared
R_squared = 1 - (RSS / TSS)
print("R-squared: %.2f" % R_squared)
```

TSS: 430.80  
 RSS: 24.62  
 R-squared: 0.94

```
In [22]: # using scikit-learn to calculate r-squared
print('R-squared: %.4f' % model.score(heights_test,
                                       weights_test))

# R-squared: 0.9429
```

R-squared: 0.9429

An R-Squared value of 0.9429 (94.29%) indicates a pretty good fit for your test data.

## Persisting the Model

```
In [23]: import pickle

# save the model to disk
filename = 'HeightsAndWeights_model.sav'
# write to the file using write and binary mode
pickle.dump(model, open(filename, 'wb'))
```

```
In [24]: # load the model from disk
loaded_model = pickle.load(open(filename, 'rb'))
```

```
In [25]: result = loaded_model.score(heights_test,
                                     weights_test)
```

Using the joblib module is very similar to using the pickle module

```
In [26]: from sklearn.externals import joblib

# save the model to disk
filename = 'HeightsAndWeights_model2.sav'
joblib.dump(model, filename)

# load the model from disk
loaded_model = joblib.load(filename)
result = loaded_model.score(heights_test,
                             weights_test)

print(result)
```

0.9428592885995253

## Data Cleansing

### Cleaning Rows with NaNs

```
In [27]: import pandas as pd
df = pd.read_csv('NaNDataset.csv')
df.isnull().sum()
```

```
Out[27]: A    0
         B    2
         C    0
         dtype: int64
```

```
In [28]: print(df)
```

	A	B	C
0	1	2.0	3
1	4	NaN	6
2	7	NaN	9
3	10	11.0	12
4	13	14.0	15
5	16	17.0	18

## Replacing NaN with the Mean of the Column

```
In [29]: # replace all the NaNs in column B with the average of column B
df.B = df.B.fillna(df.B.mean())
print(df)
```

	A	B	C
0	1	2.0	3
1	4	11.0	6
2	7	11.0	9
3	10	11.0	12
4	13	14.0	15
5	16	17.0	18

## Removing Rows

```
In [30]: df = pd.read_csv('NaNDataset.csv')
df = df.dropna()                                # drop all rows with NaN
print(df)
```

	A	B	C
0	1	2.0	3
3	10	11.0	12
4	13	14.0	15
5	16	17.0	18

```
In [31]: df = df.reset_index(drop=True)        # reset the index
print(df)
```

	A	B	C
0	1	2.0	3
1	10	11.0	12
2	13	14.0	15
3	16	17.0	18

## Removing Duplicate Rows

```
In [32]: import pandas as pd
df = pd.read_csv('DuplicateRows.csv')
print(df.duplicated(keep=False))
```

```
0    False
1     True
2     True
3    False
4    False
5     True
6     True
7    False
8    False
dtype: bool
```

```
In [33]: print(df.duplicated(keep="first"))
```

```
0    False
1    False
2     True
3    False
4    False
5    False
6     True
7    False
8    False
dtype: bool
```

```
In [34]: print(df[df.duplicated(keep=False)])
```

```
   A  B  C
1  4  5  6
2  4  5  6
5 10 11 12
6 10 11 12
```

```
In [35]: df.drop_duplicates(keep='first', inplace=True) # remove duplicates and keep first
print(df)
```

```
   A  B  C
0  1  2  3
1  4  5  6
3  7  8  9
4  7 18  9
5 10 11 12
7 13 14 15
8 16 17 18
```

```
In [36]: df.drop_duplicates(subset=['A', 'C'], keep='last',
                           inplace=True)      # remove all duplicates in
                                              # columns A and C and keep
                                              # the last

print(df)
```

	A	B	C
0	1	2	3
1	4	5	6
4	7	18	9
5	10	11	12
7	13	14	15
8	16	17	18

## Normalizing Columns

```
In [37]: import pandas as pd
from sklearn import preprocessing

df = pd.read_csv('NormalizeColumns.csv')
print(df)

x = df.values.astype(float)

min_max_scaler = preprocessing.MinMaxScaler()
x_scaled = min_max_scaler.fit_transform(x)
df = pd.DataFrame(x_scaled, columns=df.columns)
print(df)
```

	A	B	C
0	1000	2	3
1	400	5	6
2	700	6	9
3	100	11	12
4	1300	14	15
5	1600	17	18

	A	B	C
0	0.6	0.000000	0.0
1	0.2	0.200000	0.2
2	0.4	0.266667	0.4
3	0.0	0.600000	0.6
4	0.8	0.800000	0.8
5	1.0	1.000000	1.0

## Removing Outliers

### Tukey Fences

```
In [38]: import numpy as np

def outliers_iqr(data):
    q1, q3 = np.percentile(data, [25, 75])
    iqr = q3 - q1
    lower_bound = q1 - (iqr * 1.5)
    upper_bound = q3 + (iqr * 1.5)
    return np.where((data > upper_bound) | (data < lower_bound))
```

```
In [39]: import pandas as pd
df = pd.read_csv("http://www.mosaic-web.org/go/datasets/galton.csv")
print(df.head())
```

	family	father	mother	sex	height	nkids
0	1	78.5	67.0	M	73.2	4
1	1	78.5	67.0	F	69.2	4
2	1	78.5	67.0	F	69.0	4
3	1	78.5	67.0	F	69.0	4
4	2	75.5	66.5	M	73.5	4

```
In [40]: print("Outliers using outliers_iqr()")
print("=====")
for i in outliers_iqr(df.height)[0]:
    print(df[i:i+1])
```

```
Outliers using outliers_iqr()
=====
      family  father  mother  sex  height  nkids
288      72    70.0    65.0   M    79.0      7
```

## Z-Score

```
In [41]: def outliers_z_score(data):
    threshold = 3
    mean = np.mean(data)
    std = np.std(data)
    z_scores = [(y - mean) / std for y in data]
    return np.where(np.abs(z_scores) > threshold)
```



```
In [42]: print("Outliers using outliers_z_score()")
print("=====")
for i in outliers_z_score(df.height)[0]:
    print(df[i:i+1])
print()
```

Outliers using outliers\_z\_score()

=====

	family	father	mother	sex	height	nkids
125	35	71.0	69.0	M	78.0	5
	family	father	mother	sex	height	nkids
288	72	70.0	65.0	M	79.0	7
	family	father	mother	sex	height	nkids
672	155	68.0	60.0	F	56.0	7

```
In [ ]:
```